


PACKET

Ana María Matas · Week 9 Business Bending
Chapter 8 - The Job Savior
Working slice: First-Proof Evidence Capture



Capturing one human + AI behavior: noticing missing information before acting.

This packet defines one working slice of the team's First-Proof vacuum. The prototype does not certify job readiness. It records one narrow behavior under known conditions and makes the limits of that evidence visible.

1. Problem in my words

The first-rung problem is not only that young people lack formal experience. Many already do useful work in school projects, family businesses, informal jobs, internships or other settings, but that capability often does not become portable evidence that an employer can inspect and trust.

AI makes the signaling problem harder because polished output is cheaper to produce. A clean spreadsheet, proposal, summary or analysis no longer tells us whether the person understood the task, noticed missing information or simply accepted a plausible AI answer.

My TECHNOLOGIST slice focuses on one behavior that a polished final artifact can hide: **whether the person notices that critical information is missing before acting**. The product captures what information was available, what the AI produced anyway, what the participant noticed, what they did next and what the resulting evidence can and cannot support.

2. Exact user

Primary user: a university or technical-school student in Mexico seeking a first formal administrative, operations or analyst-type role, with little or no formal work history in that field.

The prototype user is comfortable enough with common digital tools and AI to request help, but may over-trust a polished answer when it sounds confident. They may already have experience from school, freelance, family-business or informal work; this task is not meant to erase that history. It is only used when a specific evidence gap still matters.

Secondary user: an employer, supervisor or approved reviewer who later inspects the First-Proof record as one input to a next-step decision, such as advance to interview, request more evidence or do not advance.


3. Success definition

Before the module closes, a user can complete one simulated work case with incomplete information, interact with a simulated AI assistant, make a decision, and generate a First-Proof evidence record that clearly shows whether they noticed the missing information before acting, including what the evidence does and does not prove.

Success is not “the candidate is job-ready.” Success is a narrow, inspectable evidence record tied to one observed behavior.


4. Image-generated mockup

The image-generated mockup below shows the key user states for this slice: task intake with incomplete information, simulated AI output, the participant decision, and the resulting evidence record. All people, tasks and data are invented for the prototype.



5. Feature flow - rendered Mermaid

The flow is intentionally simple. The system exposes one bounded task, preserves a critical information gap, allows AI assistance, records the participant's choice, and then generates evidence that states both the observed behavior and the limits of inference.




5A. Feature flow - Mermaid source

This is the Mermaid source saved in the packet so the same logic can render inside GitHub.

```
graph TD
    A[Participant opens simulated work task] --> B[System loads structured case data]
    B --> C{Is all critical information present?}
    C -->|No| D[Critical field remains missing]
    C -->|Yes| E[Participant can complete task normally]
    D --> F[Participant may request AI assistance]
    F --> G[AI produces plausible draft using available data]
    G --> H{Participant decision}
    H -->|Continue anyway| I[Record: gap not addressed]
    H -->|Flag missing info| J[Record: gap recognized]
    H -->|Ask human| K[Record: gap recognized + human dependency used]
    I --> L[Generate First-Proof evidence record]
    J --> L
    K --> L
    L --> M[Show observed behavior, AI use, missing info, revisions, dependencies]
    M --> N[Show what evidence supports]
    N --> O[Show what evidence does NOT prove]
    O --> P[Employer / reviewer inspects evidence]
```

6. Actor swimlane - rendered Mermaid + source

The swimlane makes the boundary between actors explicit. The candidate makes the decision, the AI can produce a plausible draft, the system records evidence, and a human reviewer owns the next hiring-related decision.



Mermaid source

```
graph LR
    subgraph CANDIDATE ["Candidate"]
        C1[Open task] --> C2[Review data]
        C2 --> C3[Request AI help]
        C3 --> C4[Choose: continue / flag / ask human]
    end

    subgraph SYSTEM ["System (First-Proof)"]
        S1[Load structured case] --> S2[Mark critical missing field]
        S3[Log candidate action] --> S4[Create evidence record]
        S4 --> S5[Display limits of inference]
    end

    subgraph AI ["AI (simulated)"]
        A1[Generate plausible draft from incomplete data]
    end

    subgraph REVIEWER ["Human reviewer / employer"]
        H1[Receive clarification request]
        H2[Inspect evidence record] --> H3[Use evidence for next decision]
    end

    C1 --> S1
    C2 --> S2
    C3 --> A1
    A1 --> C4
    C4 -- "if ask human" --> H1
    S2 --> S3
    S3 --> S4
    S4 --> S5
    S5 --> H2
    H2 --> H3
```

7. Benchmark and research line

Best existing solution on Earth: Parker Dewey and Riipen already demonstrate that bounded employer projects can create observable work evidence and employer feedback.

Mine differs / localizes by: instead of treating project completion or a polished final output as the proof, this slice captures whether the participant recognizes missing information before acting with AI and records the limits of what that evidence can honestly prove in a Mexico-focused first-rung context.

The prototype therefore does not claim novelty in short projects, AI assistance or employer feedback. The narrower experiment is whether evidence of judgment under incomplete information can be captured in a way that is more meaningful than a generic completion badge.

8. Long view - three years

If this slice works, First-Proof becomes a portable evidence infrastructure that accepts existing real-world evidence first and uses targeted Micro-Experiences only when an important evidence gap remains. Each proof object preserves provenance, AI assistance, observation conditions, human dependencies, uncertainty and the limits of inference instead of collapsing everything into an employability score. Over three years, universities, employers and training programs could use the same evidence format to support real hiring decisions without forcing workers to erase informal experience or pay to become legible.

9. Scope cut

- No CV builder or job board.
- No candidate ranking system or employability score.
- No full hiring platform or marketplace.
- No real employer integrations or payment flows.
- No real worker verification in this prototype.
- No automatic upgrade from DECLARED to VOUCHED or VERIFIED.
- No predictive claim that this behavior guarantees future job performance.
- No multiple industries or large case library.
- No real personal data.
- One simulated evidence-gap task only.

10. Architecture + Dragon Stack

Layer	Prototype	Purpose
Front end	Next.js + Tailwind	Task flow and evidence viewer
LLM	OpenAI API or simulated fallback	Produce a plausible draft from incomplete information
Structured / verified data	Typed local JSON / case schema	Define known fields, the missing field and allowed actions
Third Dragon component	Deterministic rule / matching logic	Map participant action to structured evidence
Persistence	Local state for MVP	Hold one simulated attempt without personal data
Deployment	Vercel free tier	Live URL
Version control	GitHub	5+ commits, 2 deploys, visible build history

AI boundary

The LLM may generate a plausible draft using the information available in the case. It does not decide whether the participant is competent, job-ready or deserving of an opportunity. The evidence claim comes from the recorded interaction and structured case state, not from an AI score.

Main simulated case

Scenario: prepare a short proposal for a client event using the documents provided. One or more decision-critical details are intentionally missing. The AI still produces a polished draft. The observed behavior is whether the participant notices that information is missing before finalizing the work.

Observable claim: “In this task, the participant recognized that critical information was missing before acting.”

11. Blueprint conditions + shadow clause

Condition 1 - provenance and scope: the evidence record must show what was observed, the relevant AI assistance, missing information, revisions, human dependencies and what cannot honestly be inferred. A verified artifact is not automatically a verified skill.

Condition 2 - existing work first: the full product considers prior evidence before creating a new task. This prototype represents the targeted Micro-Experience used only when a specific evidence gap needs fresh observation.

Condition 3 - bounded and reviewable: this classroom task has fixed scope and does not create productive unpaid work for a real business.

Condition 4 - tied to a real decision: the evidence is designed to be inspectable by a reviewer as one input to a hiring-related next step.

Condition 5 - worker does not pay: participation in the prototype is free to the worker.

Condition 6 - SHADOW CLAUSE: the full system may not repair the first rung by splitting recurring junior work into disposable Micro-Experiences. Repeated similar tasks, repeated cadence or growing cumulative hours must trigger human review and movement toward a real internship or junior role.

12. Security Floor

- No secrets in code or repo; API keys only in Vercel environment variables or local gitignored files.
- No real personal data in demo or seeds; all identities and task facts are simulated.
- If user persistence is later added, authentication is required.
- If a Supabase user table is introduced, Row Level Security must be enabled.
- Every form validates length, type and allowed values.
- Every AI-generated response is visibly labeled SIMULATED AI OUTPUT.

13. Mechanical test plan

#	Test	Expected result
1	Open the live case	SIMULATED TASK is visible and case loads correctly.
2	Request AI assistance	Output is labeled SIMULATED AI OUTPUT.
3	Continue without flagging gap	Evidence records that the gap was not addressed.
4	Flag missing information	Evidence records gapRecognized = true.
5	Ask a human reviewer	Evidence records humanDependencyUsed = true.
6	Inspect evidence record	Supports claim and "does not prove" section are both visible.
7	Try malformed / overlong input	Validation blocks or safely handles invalid input.
8	Security check	No API key, real personal data or raw prompt is exposed.
9	Bug cycle	Find one real bug, fix it, commit and redeploy.

14. Persona test

Synthetic user: Mariana, 22, State of Mexico, finishing a business / industrial-engineering-type degree. She has helped with administrative work in a family business, uses ChatGPT for school and spreadsheets, but tends to trust polished AI answers when they sound confident. She wants her first formal analyst or operations role.

In a fresh chat, show Mariana each product screen in order and ask her to attempt the task while narrating hesitation, assumptions and what she thinks the evidence record means. Log every confusion. The highest-risk usability hypothesis is that the polished AI answer makes her overlook the missing information or that she interprets the final proof as a general employability judgment.

15. Acceptance criteria

- One simulated task works end to end.
- The case contains a deliberate critical information gap.
- AI assistance can be requested and is visibly labeled simulated.
- The participant can continue, flag the gap or ask a human.
- The selected action is captured correctly.
- The evidence record states what was observed and what cannot be inferred.
- No employability score or automatic hiring recommendation appears.
- No real personal data appears.
- At least one bug is found, fixed and redeployed.
- The Persona Test is documented and the single worst confusion is fixed.

16. Build and deployment discipline

- Minimum 5 meaningful commits.
- Deployment 1 after the end-to-end task and evidence record work.
- Run the mechanical pass after Deployment 1.
- Document and fix one real bug.
- Run the Persona Test and fix the single worst usability problem.
- Deployment 2 after the bug/persona fix.
- Every session closes with DECISIONS.md updated, tomorrow's first move recorded, commit and push.